

Solutions for Homework 1

1.

(a) P1 is irreducible and aperiodic. The chain has a unique stationary distribution that it always converges to.

P2 is not irreducible. P2 is aperiodic. The chain does not have a unique stationary distribution. Hence it is also impossible for the chain to always converge to a unique distribution.

P3 is irreducible but not aperiodic. The chain has a unique stationary distribution. But it does not always converge to this stationary distribution.

(b) Since Ax_k is known, then x_{k+1} will have a Gaussian distribution $\mathcal{N}(Ax_k, W)$. Next, we can compute the mean and covariance of x_{k+1} as

$$\begin{aligned}\mathbb{E}x_{k+1} &= A\mathbb{E}x_k \\ \mathbb{E}x_{k+1}x_{k+1}^\top &= A(\mathbb{E}x_kx_k^\top)A^\top + W\end{aligned}$$

2

(a) Denote the policy in the problem statement as μ . We can solve the Bellman equation. The state value function J_μ is just a vector. We can solve J_μ as $J_\mu = (I - \gamma P_\mu)^{-1} \bar{c}_\mu$. Here $\bar{c}_\mu(i) = 0.6c(i, a_1) + 0.4c(i, a_2)$. The (i, j) -th entry of P_μ is defined as $0.6P(j, i, a_1) + 0.4P(j, i, a_2)$ where $P(j, i, a) := P(s_{t+1} = j | s_t = i, a_t = a)$.

We can also solve the Bellman equation for the Q -factor. Specifically, we have

$$Q(i, a) = c(s, a) + \gamma \sum_{j=1}^n P(j, i, a) [0.6Q(j, a_1) + 0.4Q(j, a_2)]$$

which is equivalent to another linear equation $Q_\mu = \hat{c}_\mu + \gamma M_\mu Q_\mu$ where the i -th row of M_μ is $[0.6P(1, i, a_1) \ 0.4P(1, i, a_2) \ 0.6P(2, i, a_1) \ 0.4P(2, i, a_2) \ \dots \ 0.6P(n, i, a_1) \ 0.4P(n, i, a_2)]$, and (Q_μ, \hat{c}_μ) can be calculated as

$$Q_\mu = \begin{bmatrix} Q(1, a_1) \\ Q(1, a_2) \\ Q(2, a_1) \\ Q(2, a_2) \\ \vdots \\ Q(n, a_1) \\ Q(n, a_2) \end{bmatrix}, \quad \hat{c}_\mu = \begin{bmatrix} c(1, a_1) \\ c(1, a_2) \\ c(2, a_1) \\ c(2, a_2) \\ \vdots \\ c(n, a_1) \\ c(n, a_2) \end{bmatrix}$$

Then Q_μ can be calculated as $Q_\mu = (I - \gamma M_\mu)^{-1} \hat{c}_\mu$. When the transition model is unknown, one can apply Monte Carlo simulation or temporal difference learning to learn value functions directly.

(b) At every t , the action a_t is generated using the policy μ given in Problem 2(a). Then apply a_t and measure s_{t+1} and $c(s_t, a_t)$. Then update the Q -factor as

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha_t (c(s_t, a_t) + \gamma \max[Q_t(s_{t+1}, a_1), Q_t(s_{t+1}, a_2)] - Q_t(s_t, a_t))$$

The size of the Q -table is $2n$. If $s_t = i$ and $a_t = a$, then we only update the $(2i + a - 1)$ -th entry of the Q -table at step t .

(c) For SARSA, we need to specify an initial action a_0 (which can be generated arbitrarily). At every step t , apply the action a_t , and measure s_{t+1} and $c(s_t, a_t)$. Use Q_t to generate an ε -greedy policy and then use this policy to sample an action a_{t+1} . Then update the Q -factor as

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha_t (c(s_t, a_t) + \gamma Q_t(s_{t+1}, a_{t+1}) - Q_t(s_t, a_t))$$

So at step $t \geq 1$, the action a_t is already generated using the ε -greedy policy based on Q_{t-1} .

We can see that Q -learning is off-policy in the sense that the choice of behavior policy can be independent of Q_t . In contrast, SARSA is on policy since the behavior policy is directly related to Q_t . Another difference is that in the update rules, Q -learning requires calculating $\max_{a'} Q_t(s_{t+1}, a')$ (which is equal to $\max[Q_t(s_{t+1}, a_1), Q_t(s_{t+1}, a_2)]$ in this problem) and SARSA directly applies $Q_t(s_{t+1}, a_t)$.

(d) Check Pages 3-5 of the pdf file at the following link:

<https://uofi.app.box.com/s/sniit2g18p41rgdbmb2ccvkouy7z9hkv>

3

(a) Given a linear policy K , it is straightforward to use induction to show

$$V(x) = r_K + x^\top \left(\sum_{t=0}^{\infty} \gamma^t ((A - BK)^\top)^t (Q + K^\top R K) (A - BK)^t \right) x \quad (1)$$

where r_K is some extra term introduced by the noise w_t . Therefore, we can parameterize the value function as $x^\top P_K x + r_K$. Therefore, we have

$$V(x) = x^\top (Q + K^\top R K) x + \gamma (\mathbb{E}((A - BK)x + w)^\top P_K ((A - BK)x + w) + r_K) \quad (2)$$

Notice w is independent from x and has a zero mean, we have

$$\mathbb{E}((A - BK)x + w)^\top P_K ((A - BK)x + w) = x^\top (A - BK)^\top P_K (A - BK) x + \mathbb{E}(w^\top P_K w)$$

Notice that the left side of (2) is just $x^\top P_K x + r_K$. Hence (2) can be rewritten as

$$x^\top P_K x + r_K = x^\top (Q + K^\top R K) x + \gamma x^\top (A - BK)^\top P_K (A - BK) x + \gamma \mathbb{E}(w^\top P_K w) + \gamma r_K$$

To ensure that the quadratic functions on the left and right sides of the above equation are the same, the following have to be true:

$$\begin{aligned} x^\top P_K x &= x^\top (Q + K^\top R K) x + \gamma x^\top (A - BK)^\top P_K (A - BK) x \\ r_K &= \gamma \mathbb{E}(w^\top P_K w) + \gamma r_K \end{aligned}$$

Hence, the Bellman equation becomes

$$P_K = Q + K^\top R K + \gamma (A - BK)^\top P_K (A - BK)$$

and $r_K = \frac{\gamma}{1-\gamma} \mathbb{E}(w^\top P_K w) = \frac{\gamma}{1-\gamma} \text{trace}(PW)$ where W is the covariance matrix of w_t .

For the \mathcal{Q} -function, we have

$$\begin{aligned} \mathcal{Q}(x, u) &= x^\top Q x + u^\top R u + \gamma \mathbb{E}V(Ax + Bu + w) \\ &= x^\top Q x + u^\top R u + \gamma \mathbb{E}(Ax + Bu + w)^\top P_K (Ax + Bu + w) + \gamma r_K \\ &= x^\top Q x + u^\top R u + \gamma (Ax + Bu)^\top P_K (Ax + Bu) + \gamma (\mathbb{E}(w^\top P_K w) + r_K) \\ &= \begin{bmatrix} x \\ u \end{bmatrix}^\top \begin{bmatrix} Q + \gamma A^\top P_K A & \gamma A^\top P_K B \\ \gamma B^\top P_K A & R + \gamma B^\top P_K B \end{bmatrix} \begin{bmatrix} x \\ u \end{bmatrix} + r_K \end{aligned}$$

We can also directly parameterize $\mathcal{Q}(x, u)$ as

$$\mathcal{Q}(x, u) = \begin{bmatrix} x \\ u \end{bmatrix}^\top \begin{bmatrix} \mathcal{Q}_{11} & \mathcal{Q}_{12} \\ \mathcal{Q}_{12}^\top & \mathcal{Q}_{22} \end{bmatrix} \begin{bmatrix} x \\ u \end{bmatrix} + r_K$$

Notice $V(x) = \mathcal{Q}(x, -Kx)$. Therefore, we can substitute this into the above equation to obtain the Bellman equation for \mathcal{Q} :

$$\mathcal{Q}(x, u) = x^\top Q x + u^\top R u + \gamma \mathbb{E} \mathcal{Q}(Ax + Bu + w, -K(Ax + Bu + w))$$

which is equivalent to

$$\begin{aligned} \begin{bmatrix} \mathcal{Q}_{11} & \mathcal{Q}_{12} \\ \mathcal{Q}_{12}^\top & \mathcal{Q}_{22} \end{bmatrix} &= \begin{bmatrix} Q & 0 \\ 0 & R \end{bmatrix} + \gamma \begin{bmatrix} A & B \\ -KA & -KB \end{bmatrix}^\top \begin{bmatrix} \mathcal{Q}_{11} & \mathcal{Q}_{12} \\ \mathcal{Q}_{12}^\top & \mathcal{Q}_{22} \end{bmatrix} \begin{bmatrix} A & B \\ -KA & -KB \end{bmatrix} \\ r_K &= \gamma \mathbb{E} \begin{bmatrix} w \\ -Kw \end{bmatrix}^\top \begin{bmatrix} \mathcal{Q}_{11} & \mathcal{Q}_{12} \\ \mathcal{Q}_{12}^\top & \mathcal{Q}_{22} \end{bmatrix} \begin{bmatrix} w \\ -Kw \end{bmatrix} + \gamma r_K \end{aligned}$$

(b) Optimal Bellman equation: Suppose the optimal state value function is $x^\top P x + r$. We have

$$\begin{aligned} x^\top P x + r &= \min_u (x^\top Q x + u^\top R u + \gamma \mathbb{E}(Ax + Bu + w)^\top P (Ax + Bu + w) + \gamma r) \\ &= \min_u (x^\top Q x + u^\top R u + \gamma (Ax + Bu)^\top P (Ax + Bu) + \gamma \mathbb{E} w^\top P w + \gamma r) \end{aligned}$$

Taking gradient of the function on the right side with respect to u leads to

$$u = -\gamma(R + \gamma B^\top P B)^{-1} B^\top P A x.$$

which can be substituted back to get the following optimal Bellman equation:

$$P = Q + \gamma A^\top P A - \gamma^2 A^\top P B (R + \gamma B^\top P B)^{-1} B^\top P A$$

Then the optimal state-action value function can be calculated as

$$\mathcal{Q}^*(x, u) = \begin{bmatrix} x \\ u \end{bmatrix}^\top \begin{bmatrix} Q + \gamma A^\top P A & \gamma A^\top P B \\ \gamma B^\top P A & R + \gamma B^\top P B \end{bmatrix} \begin{bmatrix} x \\ u \end{bmatrix} + \frac{\gamma}{1 - \gamma} \text{trace}(P W)$$

(c) Policy iteration: The PI algorithm iterates as $K^{n+1} = \gamma(\gamma B^\top P^n B + R)^{-1} B^\top P^n A$ where P_n solves the Bellman equation $\gamma(A - B K^n)^\top P^n (A - B K^n) + Q + (K^n)^\top R K^n = P^n$. Another option is to evaluate \mathcal{Q} for every step and then design a policy which is the greedy policy for \mathcal{Q} . Specifically, at every step n , we first solve the \mathcal{Q} Bellman equation to obtain $(\mathcal{Q}_{11}^n, \mathcal{Q}_{12}^n, \mathcal{Q}_{22}^n)$ (the policy evaluation step), and then update the policy as $K_{n+1} = -(\mathcal{Q}_{22}^n)^{-1}(\mathcal{Q}_{12}^n)^\top$ (the policy improvement step).

To estimate \mathcal{Q} -Factor from data, one can either use Monte Carlo simulation or LSTDQ (see the LQR note for the details).

(d) For SARSA, the initial action u_0 can be arbitrary. At every step n , apply the control action u_n and measure x_{n+1} and $c(x_n, u_n) = x_n^\top Q x_n + u_n^\top R u_n$. Choose u_{n+1} using the ε -greedy policy generated by $\mathcal{Q}_n(x, u) = \theta_n^\top \phi(x, u)$ where ϕ is the feature. Then update the weight vector as $\theta_{n+1} = \theta_n + \alpha_n \phi(x_n, u_n)(c(x_n, u_n) + \gamma \theta_n^\top \phi(x_{n+1}, u_{n+1}) - \theta_n^\top \phi(x_n, u_n))$. For every $n \geq 1$, u_n was sampled using the ε -greedy policy generated by $\theta_{n-1}^\top \phi(x, u)$. For Q -learning, we can choose any behavior policy that provides sufficient exploration. At every step n , sample u_n using the behavior policy and measure x_{n+1} and $c(x_n, u_n)$. Then update the weight vector as $\theta_{n+1} = \theta_n + \alpha_n \phi(x_n, u_n)(c(x_n, u_n) + \gamma \min_u \theta_n^\top \phi(x_{n+1}, u) - \theta_n^\top \phi(x_n, u_n))$. Again, Q -learning is off-policy and the sampling can be done using any behavior policy providing sufficient exploration. SARSA is on-policy and sampling is done using the ε -greedy policy given by $\theta_n^\top \phi(x, u)$.

4

(a) Popular options for Ψ_t :

- Monte Carlo estimation: $\sum_{t'=t}^{\infty} \gamma^{t'-t} c_{t'}$
- Baselined versions of Monte Carlo estimation: $\sum_{t'=t}^{\infty} (\gamma^{t'-t} c_{t'} - b(s_t))$
- State-action value function: $Q^\pi(x_t, u_t)$

- Advantage function: $A^\pi(x_t, u_t)$
- TD residual: $c_t + \gamma V^\pi(x_{t+1}) - V^\pi(x_t)$
- Generalized advantage estimation

To calculate the gradient, first notice K is a matrix. Hence $\nabla_\theta \log \pi_\theta$ is also a matrix. The (i, j) -th entry of this matrix is just

$$\frac{\partial \log \pi_\theta}{\partial K_{ij}} = -\sigma^{-1}(u_t^{(i)} + \sum_{j=1}^{n_x} K_{ij} x_t^{(j)}) x_t^{(j)}$$

More compactly, we can write $\nabla_\theta \log \pi_\theta(u_t | x_t) = -(\sigma I)^{-1}(u_t + K x_t) x_t^\top$.

Now consider the case $u_t \sim \mathcal{N}(W^1 h(W^0 x_t), \sigma I)$. The derivative with respect to W^1 can be directly calculated as

$$\frac{\partial}{\partial W^1} \log \pi_\theta(u_t | x_t) = \sigma^{-1}(u_t + W^1 h(W^0 x_t))(h(W^0 x_t))^\top$$

The derivative with respect to W^0 requires a backpropagation step and can be calculated as

$$\frac{\partial}{\partial W^0} \log \pi_\theta(u_t | x_t) = (W^1 \text{diag}(h'(W^0 x_t)))^\top (u_t - W^1 h(W^0 x_t)) x_t^\top$$

where $\text{diag}(h'(W^0 x_t))$ is a diagonal matrix whose (i, i) -th entry is equal to the i -th entry of the vector $h'(W^0 x_t)$.

(b)

$$\begin{aligned} \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)} \left(\lim_{\sigma \rightarrow 0} \frac{\mathcal{C}(K + \sigma \varepsilon) - \mathcal{C}(K)}{\sigma} \right) \varepsilon &= \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)} (\varepsilon^\top \nabla \mathcal{C}(K)) \varepsilon \\ &= \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)} \varepsilon (\varepsilon^\top \nabla \mathcal{C}(K)) \\ &= \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)} (\varepsilon \varepsilon^\top) \nabla \mathcal{C}(K) \\ &= \nabla \mathcal{C}(K) \end{aligned}$$

5

A code will be provided on the course website. LSTD-Q works efficiently for (b), and approximate PI works efficiently for (c). See the code for how to setup behavior policy for efficient exploration.